

MET CS 622:

Concurrency

Reza Rawassizadeh

BRACE YOURSELF

**CONCURRENCY IS THE
HARDEST PART OF PROGRAMMING**

Outline

- Thread Concepts
- Locks
 - Structured Lock
 - Unstructured Locks

Outline

- **Thread Concepts**
- Locks
 - Structured Lock
 - Unstructured Locks

Problem with Disk Operation

- Performing an operation inside memory is very fast.
- In contrast, accessing disk and performing I/O is slow process.
- If a process need something from disk, usually it should wait until that particular information gets available and then uses it.
- Considering slow disk access the performance of computers could be very poor.
- Computers are multi-task and should run several things in parallel. Operating systems resolve this challenge by introducing a **multi-thread feature**, which is called **concurrency**.

Thread



Definitions

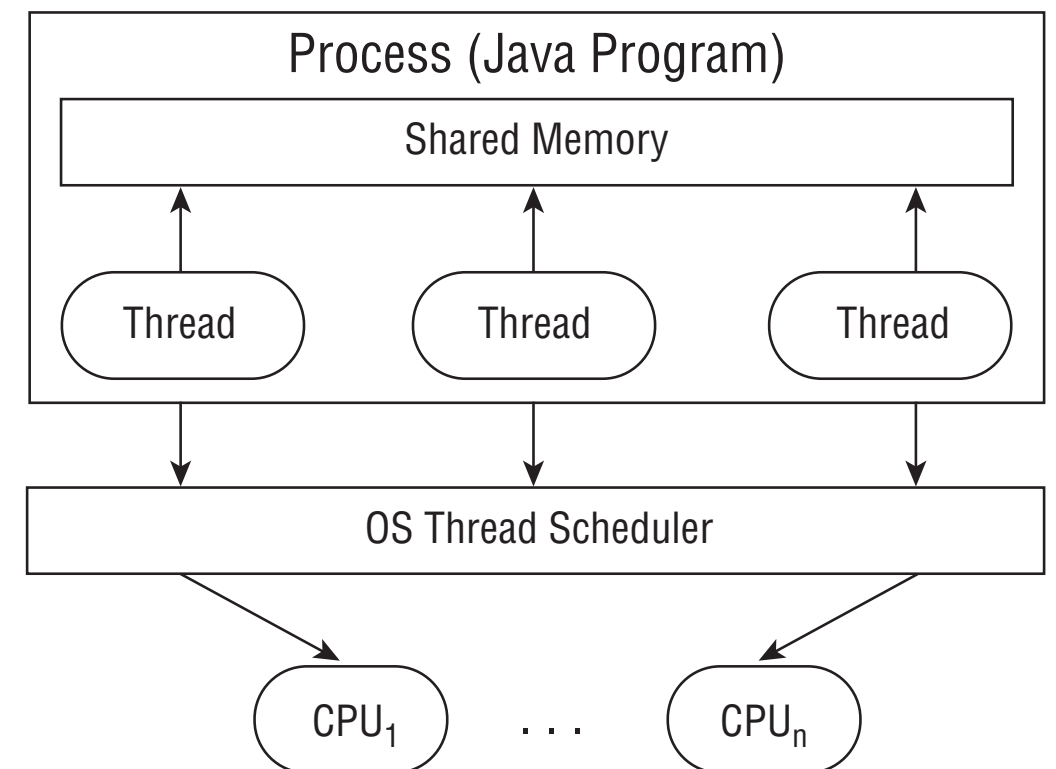
- **Thread** is the smallest unit of execution that can be scheduled by the operating system.
- A **process** is a set of threads that execute in the same shared environment (shared memory space).
- A **task** is a unit of a work that is performed by a thread.
- Executing multiple threads or processes at a time is called **concurrency**.
- Most of Java applications are multi-thread.

CPU and Thread

- A single core CPU can execute only one thread.
- A multi core CPU can execute more than one thread, but the number of threads could be larger than the number of CPU cores.
- Operating systems use something called **thread scheduler**, which decide what should be executed.

Thread Scheduler

- Operating system use thread scheduler to determine which thread should be running on a CPU core.
- Thread schedulers can use a **round-robin algorithm** to determine resource to thread assignment.
- The round-robin algorithm focus on *equal resource assignment* to each thread, in a circular fashion.
- To be honest, thread scheduling algorithms are far more complex than round-robin, but we need to be familiar with the concept of round-robin algorithm.



Source: OCP_ Oracle Certified Professional Java SE 8
Programmer II Study Guide_ Exam 1Z0-809-Sybex

Thread Priority

- Java enables us to assign a priority to a thread.
- It supports three thread priorities:
 - `Thread.MIN_PRIORITY`
 - `Thread.NORM_PRIORITY`
 - `Thread.MAX_PRIORITY`

Runnable

- `java.lang.Runnable` is a **functional java interface** that is used to define the work that a thread should execute.
- A functional interface is an interface that contains only **one abstract method**, but it can have any number of non-abstract methods.
- Until now, we use `main` as an application execution entry point. However, for thread execution we need to call the `run()` method inside the `java.lang.Runnable` interface.

```
@FunctionalInterface
public interface Runnable {
    void run();
}
```

Lambda Expression

- We can also use lambda expression that implicitly implements the Runnable interface.

`() -> System.out.println("Hello World")`

`() -> {int i=10; i++;}`

`() -> {return;}`

`() -> {}`

An example class that runs something inside a thread

```
public class ThreadExample implements Runnable {  
    public void run(){  
        // what is running inside the thread goes here  
    }  
}
```

Creating a Thread

There are two steps required to create a thread.

1. Defining the thread within the tasks it should perform.
2. Starting the thread.

There is no guarantee about the order of threads to be executed in Java.

HelloWorld Thread

```
package edu.bu.met622.threadtest;

public class SuperSimplePrint implements Runnable{

    @Override
    public void run() {
        for (int i=0; i<10; i++) {
            System.out.println("i is:"+i);
        }
    }

    public static void main(String[] args) {
        Thread t = new Thread(new SuperSimplePrint());
        t.start();
    }

}
```

Summary of Concurrency in Java

- Thread is a **class**
- Runnable is an **interface**
- The **Thread class implements Runnable**
- The argument passed to the thread constructor must be a Runnable.
- The threads spawned by a process run asynchronously

Now we make two threads and each one is printing a number (their task).

```
package edu.bu.met622.threadtest;

public class FirstThread implements Runnable{

    @Override
    public void run() {
        for (int i =0; i<100 ; i++) {
            System.out.println("-from FirstThread: i:"+i);
        }
    }
}
```

```
package edu.bu.met622.threadtest;

public class SecondThread implements Runnable{

    @Override
    public void run() {
        for (int j =0; j<100 ; j++) {
            System.out.println("-from SecondThread: j:"+j);
        }
    }
}
```

```
package edu.bu.met622.threadtest;

import edu.bu.met622.threadtest.FirstThread;
import edu.bu.met622.threadtest.SecondThread;

public class TesttwoThread {
    public static void main(String[] args) {
        FirstThread firstT = new FirstThread() ;
        Thread a = new Thread(firstT);

        SecondThread secondT = new SecondThread();
        Thread b = new Thread(secondT);

        a.start();
        b.start();
    }
}
```

Difference between `run()` and `start()`

- The method `start()`, will start a new thread and the JVM assigns it to a CPU core.
- The `run()`, will execute the content of a thread.

Some Common Job Interview Questions about Threads

- Explain the difference between extending the Thread class and implementing Runnable.
 - A. If we need to define our own thread rules, upon which multiple tasks will rely, e.g. a **priority thread**, extending thread may be preferable.
 - B. Java doesn't support multiple inheritance, and thus **extending Thread** does not allow us to extend any other class, whereas implementing Runnable lets you extend another class.
 - C. **Implementing Runnable** is often a better object-oriented design practice, because **it separates the task being performed from the Thread object that are performing it.**
 - D. Implementing Runnable allows the class to be used by numerous Concurrency API classes.

Some Common Job Interview Questions about Threads

- Explain the difference between extending the Thread class and implementing Runnable.
 - A. If we need to define our own thread rules, upon which multiple tasks will rely, e.g. a **priority thread**, extending thread may be preferable.
 - B. Java does not allow extending Thread class, but it does not allow implementing Runnable either. **Many books recommend avoid extending Thread class, unless you must doing it.**
 - C. **Implementing Runnable** is often a better object-oriented design practice, because **it separates the task being performed from the Thread object that are performing it.**
 - D. Implementing Runnable allows the class to be used by numerous Concurrency API classes.

Thread Join

experiment this code with/without join

The join method allows one thread to wait until another thread gets accepted.

```
package edu.bu.met622.threadtest;

public class TestThreadJoin {
    public static void main(String[] args) throws InterruptedException {
        Thread t1 = new Thread() {
            public void run() {
                for (int j=0 ; j<20; j++) {
                    System.out.println("t1");
                }
            }
        };
        Thread t2 = new Thread() {
            public void run() {
                for (int j=0 ; j<20; j++) {
                    System.out.println("t2");
                }
            }
        };
        t1.start();
        // Lets comment and uncomment t2.join and see the differences.
        t2.join();
        t2.start();
        System.out.println("End");
    }
}
```

Thread Sleep

- As it has been explained a threads is running out of our control after it has been started.
- Some times it is useful if we can pause thread execution temporary.
- `Thread.sleep(xxx)`, causes a thread to pauses its execution for xxx milliseconds.
- `sleep()` is a **static** method in Thread.
- The `sleep()` method throws a checked exception, namely `InterruptedException`.

More example

```
public class TestThreadJoin {
    public static void main(String[] args) throws InterruptedException {
        Thread t1 = new Thread() {
            public void run() {
                for (int j=0 ; j<100; j++) {
                    System.out.println("taaaaaa1");
                }
            }
        };
        Thread t2 = new Thread() {
            public void run() {
                for (int j=0 ; j<100; j++) {
                    System.out.println("t222222");
                }
            }
        };
        Thread t3 = new Thread() {
            public void run() {
                for (int k=0 ; k<100; k++) {
                    System.out.println("t333333");
                }
            }
        };
        t1.start();
        t1.sleep(1000);
        t2.start();
        t3.start();
        // Lets comment and uncomment t2.join and see the differences.
        // t2.join();
        System.out.println("End");
    }
}
```

Thread Sleep

experiment with/without sleep

```
package edu.bu.met622.threadtest;

public class SleepTest1 {
    public static void main(String args[]) throws InterruptedException {
        String importantInfo[] = { "msg 1", "msg 2", "msg 3", "msg 4" };

        for (int i = 0; i < importantInfo.length; i++) {
            // Pause for 1 seconds
            Thread.sleep(2000);
            // Print a message
            System.out.println(importantInfo[i]);
        }
    }
}
```

Outline

- Thread Concepts
- **Locks**
 - Structured Lock
 - Unstructured Locks


Race Condition Problem

- Suppose we have a method that count an object inside our limited memory, e.g. `increment(val)` as follows:

```
void increment(val) {  
    val = val+1  
}
```
- Now imagine we have two threads, T1 and T2, and both are calling `increment(val)`
- They are not aware about each other, but they do the same thing. When T1 calls `increment(val)` and increases `val` from zero to one, then T2 calls `val` to increments it to 1, but mistakingly it will increased to 2, because T1 did increase to 1.

Another Example

Sheep Counting



Outline

- Thread Concepts
- Locks
 - **Structured Lock**
 - Unstructured Locks

Structured Lock

- To avoid such a confusion we can use the `synchronize` keyword, which imposes a **lock** on the `increment` while a thread is working with it. Or we can use atomic classes, which will be described later.

```
synchronized void increment(val){  
    val = val+1  
}
```

- A content of the synchronized method is called **mutually exclusive**, because one thread at a time can work with it. As soon as the thread is starting to work with this method, there will be **lock** on this method and no other method can use it. When a thread is done, the lock will be released and other threads can use it.

Locking an Object with `synchronized`

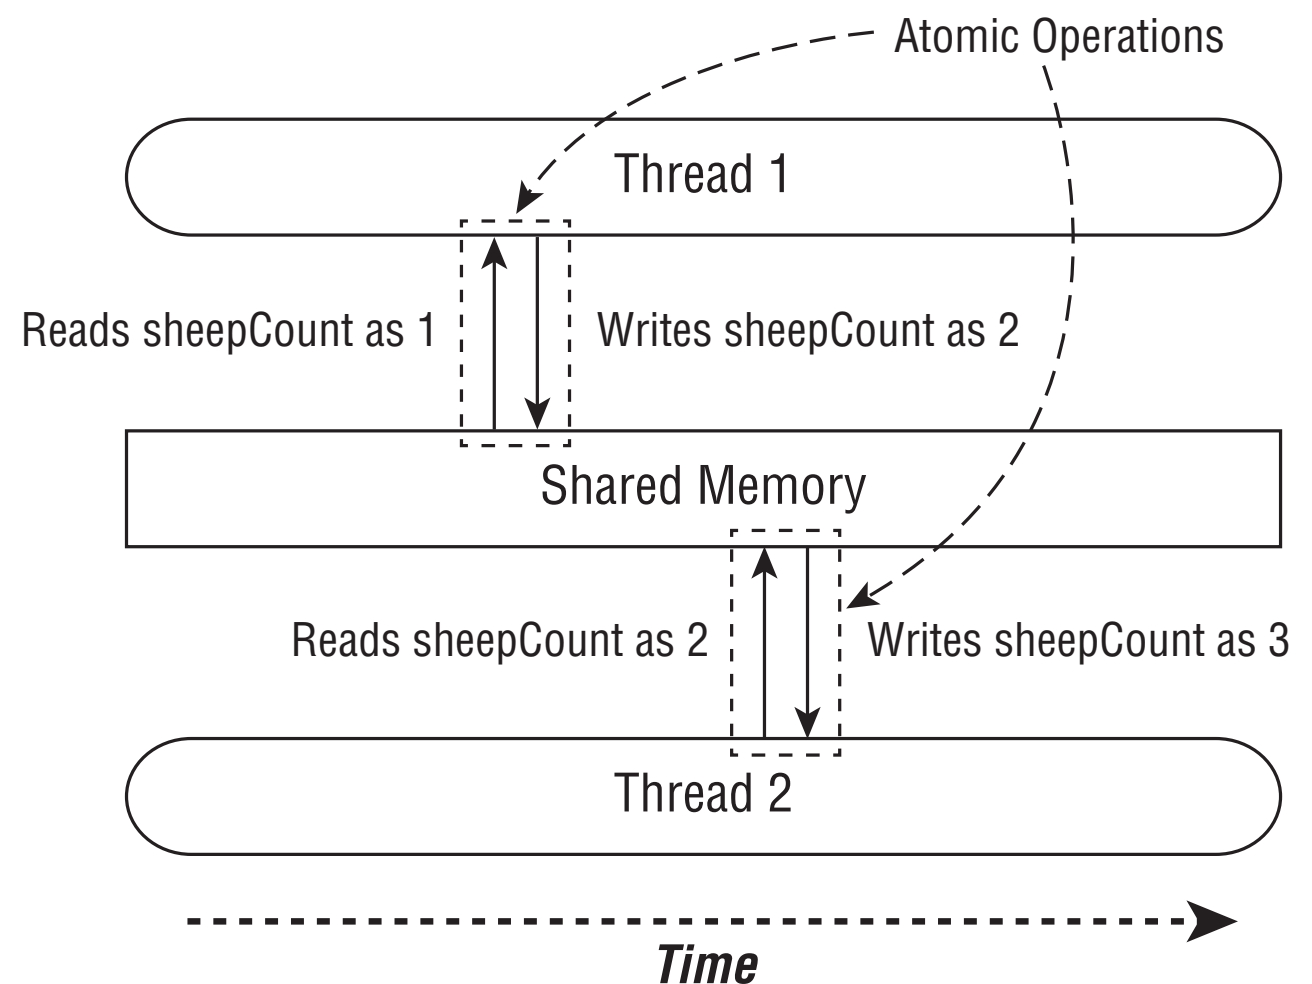
- Some time locking the entire method might negatively impact the other tasks of the method. We can also lock a an object in a block of code.

```
BufferVal buf = new BufferVal{  
    ...  
    synchronized(buf) {  
        val = val+1;  
        ...  
    }  
    ...  
}
```


Atomic

- Atomic is a property of operation, which **disables any other thread interference** on that operation.

It is more restricted than Synchronized lock. There are limited number of Atomic classes and Methods and we need to keep them in mind.



Atomic classes and Methods

- Atomic Classes: `AtomicBoolean`, `AtomicInteger`, `AtomicIntegerArray`, `AtomicLong`, `AtomicLongArray`, `AtomicReference`, `AtomicReferenceArray`
- Atomic Methods: `get()`, `set()`, `getAndSet()`, `incrementAndGet()`, `getAndIncrement()`, `decrementAndGet()`, `getAndDecrement()`

AtomicInteger Example

```
private AtomicInteger sheepCount = new AtomicInteger(0);

private void incrementAndReport() {
    System.out.print(sheepCount.incrementAndGet()+" ");
}
```

What are the disadvantage of using Synchronized for thread?

Synchronization Disadvantage

- The objective of multi-thread programming is doing task multiple tasks in parallel. By enforcing a lock we disable the multi tasking features, which might increase response time.
- Synchronization protects data integrity, but its costs will be on response time.
- Being able to identify the **performance bottleneck of a system**, especially in multithread environment is very valuable capability.

Outline

- Thread Concepts
- Locks
 - Structured Lock
 - **Unstructured Locks**

Unstructured Lock

Assume we have a four objects (a,b,c,d) and we have three works that should to be executed in a sequence, we need to do with these objects (w1, w2, w3).

Each of these works require two objects: w1(a,b), w2(b,c) , w3(c,d)




work 1(a,b)

work 2(b,c)


work 3(c,d)

Unstructured Lock


work 1(a,b) starts



work 1 done,
work 2(b,c) starts




work 2 done,
work 3(c,d) starts



work 3 done



Locks are nested inside each other



How to Resolve Nested Locks

- There are two solutions designed to handle nested locks, using a **ReentrantLock** (**Interface Lock** or **Try Lock**) and using **Read/Write Lock**.
- ReentrantLock will hold the lock until the thread is done.
- Read/Write Lock separates read and write operation and implement different locks for them

ReentrantLock

- **Lock** implementations provide more extensive locking operations than can be obtained using **synchronized** methods and statements. They allow more flexible structuring.

```
Lock lock = new ...;  
lock.lock();  
try{  
    // use the resource protected by this lock  
}  
finally {  
    lock.unlock();  
}
```

Try Lock Example (1/3)

```
package edu.bu.met622.threadtest.trylock;

public class MyThread extends Thread{

    PrintThread printT;
    MyThread(String name, PrintThread pt){
        super(name);
        this.printT = pt;
    }
    @Override
    public void run() {
        System.out.printf(
            "%s starts printing a document\n",
            Thread.currentThread().getName());
        printT.print();
    }
}
```

Try Lock Example (2/3)

```
package edu.bu.met622.threadtest.trylock;

import java.util.concurrent.locks.Lock;
import java.util.concurrent.locks.ReentrantLock;

public class PrintThread {
    private final Lock queueLock = new ReentrantLock();

    public void print() {
        queueLock.lock();

        try {
            Long duration = (long) (Math.random() * 1000);
            // Play with this duration to show lock changes
            System.out.println(Thread.currentThread().getName() +
                               " Time Taken " + (duration / 1000) + " secs.");
            Thread.sleep(duration);
        } catch (InterruptedException e) {
            e.printStackTrace();
        } finally {
            System.out.printf(
                "%s printed the document successfully.\n",
                Thread.currentThread().getName());
            queueLock.unlock();
        }
    }
}
```

Try Lock Example (3/3)

```
package edu.bu.met622.threadtest.trylock;

public class TestThread {
    public static void main(String args[]) {
        PrintThread PD = new PrintThread();

        MyThread t1 = new MyThread("Thread - 1 ", PD);
        MyThread t2 = new MyThread("Thread - 2 ", PD);
        MyThread t3 = new MyThread("Thread - 3 ", PD);
        MyThread t4 = new MyThread("Thread - 4 ", PD);

        t1.start();
        t2.start();
        t3.start();
        t4.start();
    }
}
```

Read/Write Lock

- It is a solution to mitigate nested locks.
- For example, assume we have an application that both read and write a set of objects. The read threads occurs more often than write, but some writes threads are also there.
- We can define two locks:
 - **Read Lock:** It will be unlocked if no thread is reading or writing.
 - **Write Lock:** It will be unlocked if no thread is writing and no thread has requested write access.
- *More than one thread can hold a Read Lock, but only one thread can holds the Write Lock*

```

public class ReadWriteLock{

    private int readers      = 0;
    private int writers      = 0;
    private int writeRequests = 0;

    public synchronized void lockRead() throws InterruptedException{
        while(writers > 0 || writeRequests > 0){
            wait();
        }
        readers++;
    }

    public synchronized void unlockRead(){
        readers--;
        notifyAll();
    }

    public synchronized void lockWrite() throws InterruptedException{
        writeRequests++;

        while(readers > 0 || writers > 0){
            wait();
        }
        writeRequests--;
        writers++;
    }

    public synchronized void unlockWrite() throws InterruptedException{
        writers--;
        notifyAll();
    }
}

```

source: <http://tutorials.jenkov.com/java-concurrency/read-write-locks.html>

Semantic Errors in Concurrency

(Liveness problems)

- DeadLock
- LiveLock
- Starvation

DeadLock



DeadLock

Thread 1:

```
synch(A) {  
    synch(B) {  
    }  
}
```

Thread 2:

```
synch(B) {  
    synch(A) {  
    }  
}
```

Thread 1 starts and acquires a lock on A, then Thread 2 starts and acquires a lock on B. Now thread 1 is willing to acquire the lock on B, but thread 2 is holding it. Also thread B is waiting for thread A to release the lock, but A is waiting for B.

LiveLock

Threads are not blocked, but they are in a mode, which do not make any progress toward finishing their task.

Thread 1:

```
do {  
    synch(x) {  
        x=x+1  
    }  
} while (x<2)
```

Thread 2:

```
do {  
    synch(x) {  
        x=x-1  
    }  
} while (x>-2)
```

LiveLock

Thread 1:

```
do {  
    synch(x) {  
        x=x+1  
    }  
} while (x<3)
```

Thread 2:

```
do {  
    synch(x) {  
        x=x-1  
    }  
} while (x>-3)
```

x=0

T1: x=1

T2: x=0

T1: x=1

T1: x=2

T2: x=1

T2: x=0

T2: x=-1

T1: x=0

...

Starvation

- Assume we have many threads, e.g. 500 threads (T1, ... T500) and each are designed to do a small task, which might be similar or dissimilar tasks.
- We start all of them together, and they continuously get called and execute their tasks.
- Since we have too many threads, some threads might never get called, e.g. T1,....T498, called but the last two ones never get called, T499 and T500.
- Then we can say T499 and T500 are starved.

Starvation

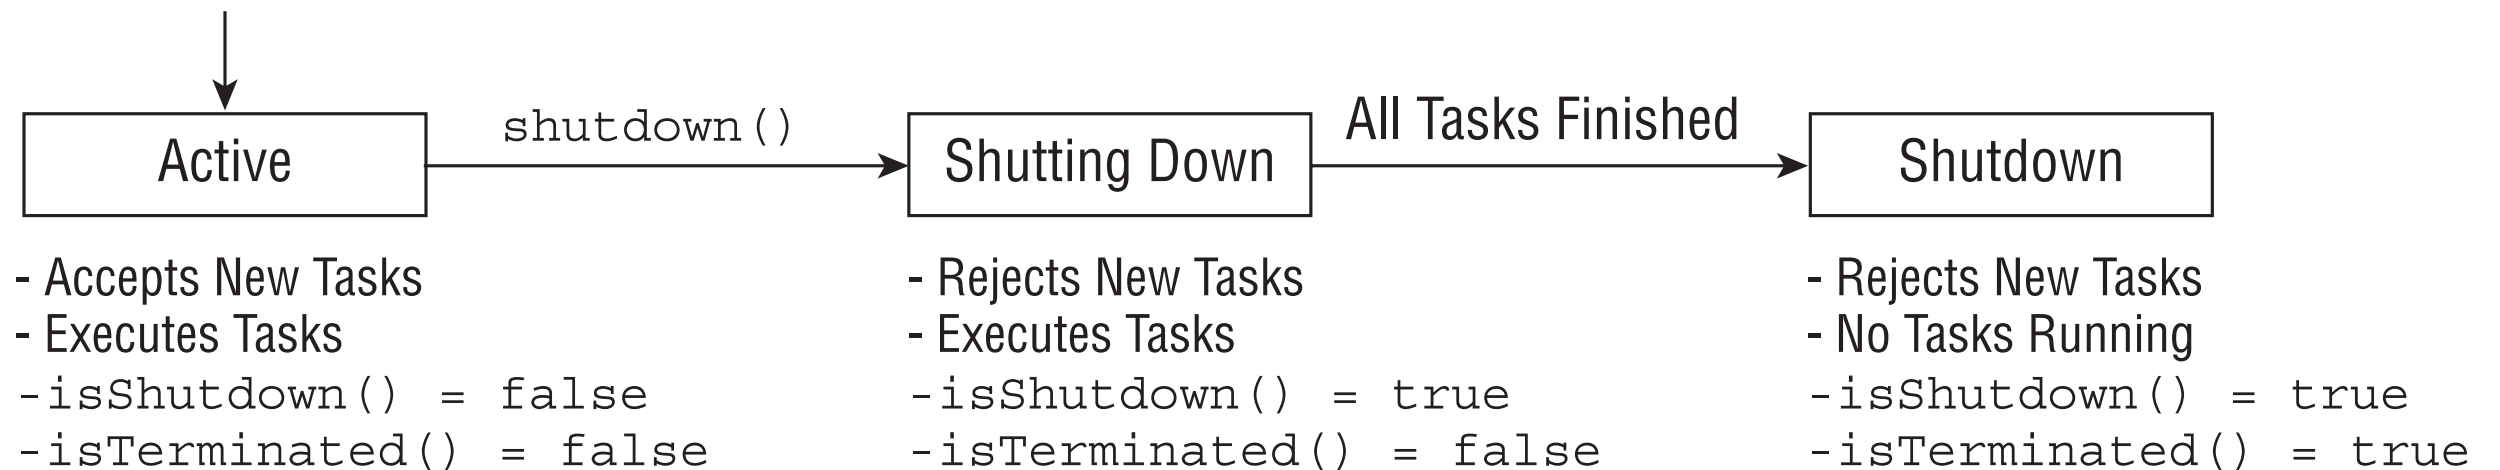
- Assume we have many threads, e.g. 500 threads (T1, ... T500) and each are designed to do a small task, which might be similar or dissimilar tasks.
- We start calling threads and they start to execute. Some threads may not get called. This usually happen in read/write locks. We do lots of read threads and write threads are getting starved.
- Since we have too many threads, some threads might never get called, e.g. T1,....T498, called but the last two ones never get called, T499 and T500.
- Then we can say T400 and T500 are starved.

ExecutorService & Thread Pools

- Previous Java thread classes are complex to make them scalable. Because thread creation and management will be handled inside other codes.
- Java tries to mitigate this challenge by introducing the **Executor Framework**. It helps us in thread creation, management and task assignment to the thread pool:
 1. **Executor**: It is an interface, with an **execute()** method to launch a task specified by a **Runnable** object.
 2. **ExecutorService**: It implements the **Executor** interface and adds functionality to manage life cycle of the **task**. It also includes **submit()** method which is a newer version of **execute()**. It returns **Callable** object. Callable objects are similar to Runnable but they have output as well.
 3. **ScheduledExecutorService**: A sub-interface of **ExecutorService**, but it has the functionality to the execution of tasks as well.

Executor LifeCycle

Create New Thread Executor



Callable vs Submit

- The `execute()` takes a `Runnable` object and completes the task. It **doesn't return anything about the success or failure of task completion** (fire-and-forget).
- The `submit()` is a newer method which **returns the future object, which is the result of thread execution**. This object can assist us identifying the result of thread execution.

`void execute(Runnable command)`: Executes a `Runnable` task at some point in the future.

`<T> Future<T> submit(Callable<T> task)`: Executes a `Callable` task at some point in the future and returns a `Future` object which is representing the pending results of the task.

Future Object

The Future class includes methods that are useful in determining the state of a task.

`boolean isDone()`: Returns true if the task was compiled.

`boolean isCancelled()`:

`boolean cancel()`: Attempt to cancel the execution of a task.

`V get()`: Retrieves the result of task.

`V get(long timeout, TimeUnit unit)`: Retrieves the result of a task, waiting the specified amount of time. If the result is not ready by the time the timeout is reached, a checked `TimeoutException` will be thrown.

Executor Example (1)

```
package edu.bu.met622.threadtest.executor2eg;

import java.util.concurrent.Callable;
import java.util.concurrent.ExecutorService;
import java.util.concurrent.Executors;

public class ExecutorExample {
    public static void main(String[] args) {
        System.out.println("Inside : " + Thread.currentThread().getName());

        // Callable (returns future object)
        System.out.println("Creating a Runnable...");
        Runnable runnable = () -> {
            // System.out.println("Inside : " + Thread.currentThread().getName());
        };

        Callable callable = Executors.callable(runnable);
        System.out.print("---Callable output:" + callable.getClass().getSimpleName());
        System.out.print("---Callable output2:" + callable.getClass().getTypeName());

        // Executor (does not return future object)
        System.out.println("Creating Executor Service...");
        ExecutorService executorService = Executors.newSingleThreadExecutor();

        // System.out.println("Submit the task specified by the runnable to the executor service.");
        executorService.submit(runnable);
        executorService.shutdown();
    }
}
```

Executor Example (2-1)

```
package edu.bu.met622.threadtest.executoreg;

class Counter {
    private int c = 0;
    public void increment() {
        System.out.println("Thread Id: " +
                           Thread.currentThread().getId());
        c = c + 1;
    }
    public int getvalue() {return c;}
}
```

Executor Example (2-2)

```
package edu.bu.met622.threadtest.executoreg;

class First implements Runnable {
    private int i;
    private Counter ctr;
    public First(int i, Counter c) {
        this.i = i;
        this.ctr = c;
    }
    public void changeCounter(Counter cntr) {
        if (i > 0) {
            synchronized(cntr) {cntr.increment();}
        }
    }
    public void run() {
        changeCounter(ctr);
    }
}
```

Executor Example (2-3)

```
package edu.bu.met622.threadtest.executoreg;

import java.util.concurrent.*;
public class ThreadPoolExecuteService
{
    public static void main(String[] args) throws InterruptedException
    {
        Counter sharedcounter = new Counter();
        ExecutorService ex = Executors.newCachedThreadPool();
        ex.execute(new First(1, sharedcounter));
        ex.execute(new First(2, sharedcounter));
        ex.execute(new First(3, sharedcounter));

        ex.shutdown();

        //      ex.awaitTermination(50, TimeUnit.MILLISECONDS);
        if (ex.isTerminated()) System.out.println("Final count = " +
sharedcounter.getvalue());
    }
}
```

References

<https://www.coursera.org/instructor/vivek-sarkar>

<https://www.callicoder.com/java-executor-service-and-thread-pool-tutorial/>

